

Time series data

```
require(ptLasso)
#> Loading required package: ptLasso
#> Loading required package: ggplot2
#> Loading required package: glmnet
#> Loading required package: Matrix
#> Loaded glmnet 4.1-8
#> Loading required package: gridExtra
```

We may have repeated measurements of X and y across time; for example, we may observe patients at two different points in time. We expect that the relationship between X and y will be different at time 1 and time 2, but not completely unrelated. Therefore, pretraining can be useful: we can use the model fitted at time 1 to inform the model for time 2.

ptLasso supports this setting, and below is an example. We first assume that X is constant across time, and y changes. Later, we will show an example where X changes across time.

To do pretraining with time series data, we:

1. fit a model for time 1 and extract its offset and support,
2. use the offset and support (the usual pretraining) to train a model for time 2.

We could continue this for k time points: after fitting a model for time 2, we would extract the offset and support. Now, the offset will include the offset from time 1 and the prediction from time 2; the support will be the *union* of supports from the first two models.

Example 1: covariates are constant over time

We'll start by simulating data – more details in the comments.

```
set.seed(1234)

# Define constants
n = 600          # Total number of samples
ntrain = 300     # Number of training samples
p = 100          # Number of features
sigma = 3        # Standard deviation of noise

# Generate covariate matrix
x = matrix(rnorm(n * p), n, p)

# Define coefficients for time points 1 and 2
beta1 = c(rep(2, 10), rep(0, p - 10)) # Coefs at time 1
beta2 = runif(p, 0.5, 2) * beta1       # Coefs at time 2, shared support with time 1

# Generate response variables for times 1 and 2
y = cbind(
  x %*% beta1 + sigma * rnorm(n),
  x %*% beta2 + sigma * rnorm(n)
)
```

```

# Split data into training and testing sets
xtest = x[-(1:ntrain), ] # Test covariates
ytest = y[-(1:ntrain), ] # Test response

x = x[1:ntrain, ] # Train covariates
y = y[1:ntrain, ] # Train response

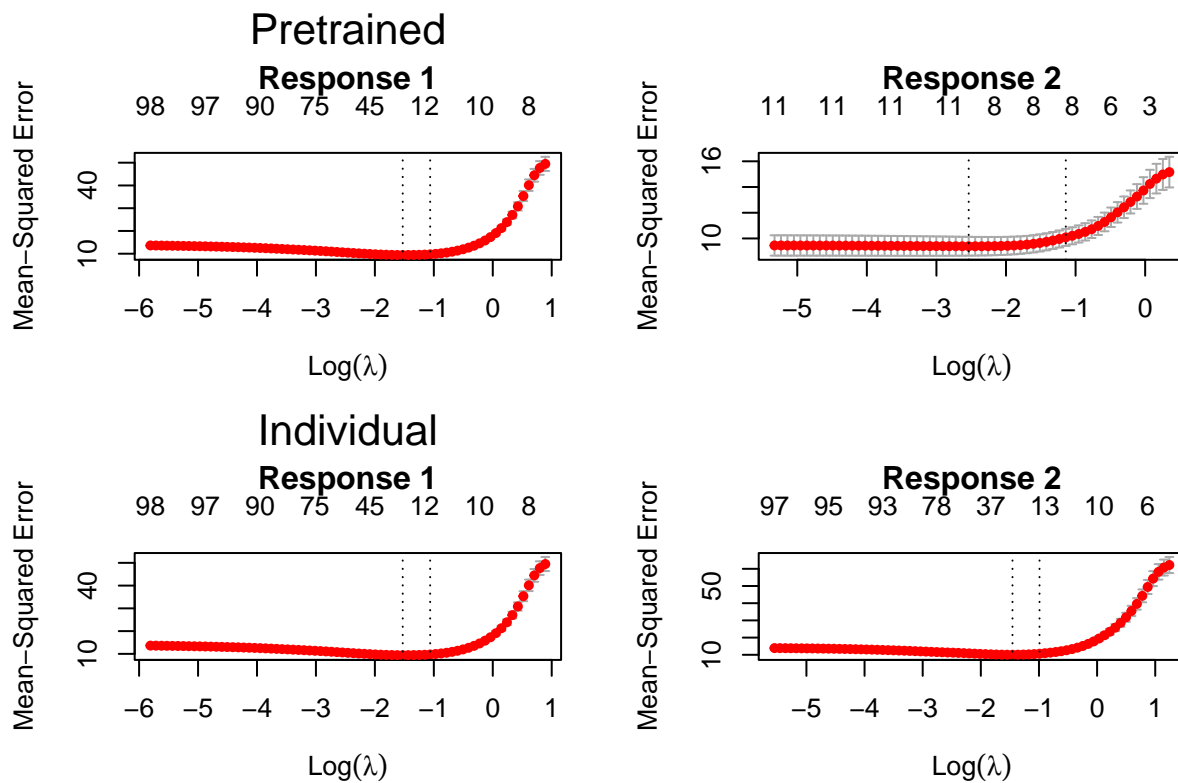
```

Having simulated data, we are ready to call `ptLasso`; the call to `ptLasso` looks much the same as in all our other examples, only now (1) `y` is a matrix with one column for each time point and (2) we specify `use.case = "timeSeries"`. After fitting, a call to `plot` shows the models fitted for both of the time points with and without using pretraining.

```

fit = ptLasso(x, y, use.case = "timeSeries", alpha = 0)
plot(fit)

```



And as before, we can `predict` with `xtest`. In this example, pretraining helps performance: the two time points share the same support, and pretraining discovers and leverages this.

```

preds = predict(fit, xtest, ytest = ytest)
preds
#>
#> Call:
#> predict.ptLasso(object = fit, xtest = xtest, ytest = ytest)
#>
#>
#>
#> alpha = 0
#>
#> Performance (Mean squared error):
#>

```

```

#>               mean response_1 response_2
#> Pretrain      9.604         10.78      8.428
#> Individual 10.428         10.78     10.076
#>
#> Support size:
#>
#> Pretrain  26 (10 common + 16 individual)
#> Individual 39

```


We specified `alpha = 0` in this example, but cross validation would advise us to choose $\alpha = 0.2$. Plotting shows us the average performance across the two time points. Importantly, at time 1, the individual model and the pretrained model are the same; we do not see the advantage of pretraining until time 2 (when we use information from time 1).

```

cvfit = cv.ptLasso(x, y, use.case = "timeSeries")
plot(cvfit)

```

Average performance over 2 responses



```

predict(cvfit, xtest, ytest = ytest)
#>
#> Call:
#> predict.cv.ptLasso(object = cvfit, xtest = xtest, ytest = ytest)
#>
#>
#>
#> alpha = 0.2
#>
#> Performance (Mean squared error):
#>

```

```
#>           mean response_1 response_2
#> Pretrain    10.62         10.87     10.37
#> Individual 10.45         10.87     10.03
#>
#> Support size:
#>
#> Pretrain    28 (10 common + 18 individual)
#> Individual  40
```

Note that we could also have treated this as a *multireponse* problem, and ignored the time-ordering of the responses. See more in the section called “Multi-response data with Gaussian responses”. (However, time ordering can be informative, and the multi-response approach does not make use of this.)

```
fit = ptLasso(x, y, use.case = "multiresponse")
```

Example 2: covariates change over time

Now, we’ll repeat what we did above, but we’ll simulate data where x changes with time. In this setting, `ptLasso` expects x to be a list with one covariate matrix for each time.

```
set.seed(1234) # Set seed for reproducibility

# Define constants
n = 600          # Total number of samples
ntrain = 300     # Number of training samples
p = 100          # Number of features
sigma = 3        # Standard deviation of noise

# Covariates for times 1 and 2
x1 = matrix(rnorm(n * p), n, p)
x2 = x1 + matrix(0.2 * rnorm(n * p), n, p) # Perturbed covariates for time 2
x = list(x1, x2)

# Define coefficients for time points 1 and 2
beta1 = c(rep(2, 10), rep(0, p - 10)) # Coefs at time 1
beta2 = runif(p, 0.5, 2) * beta1       # Coefs at time 2, shared support with time 1

# Response variables for times 1 and 2:
y = cbind(
  x[[1]] %*% beta1 + sigma * rnorm(n),
  x[[2]] %*% beta2 + sigma * rnorm(n)
)

# Split data into training and testing sets
xtest = lapply(x, function(xx) xx[-(1:ntrain), ]) # Test covariates
ytest = y[-(1:ntrain), ] # Test response

x = lapply(x, function(xx) xx[1:ntrain, ]) # Train covariates
y = y[1:ntrain, ] # Train response
```

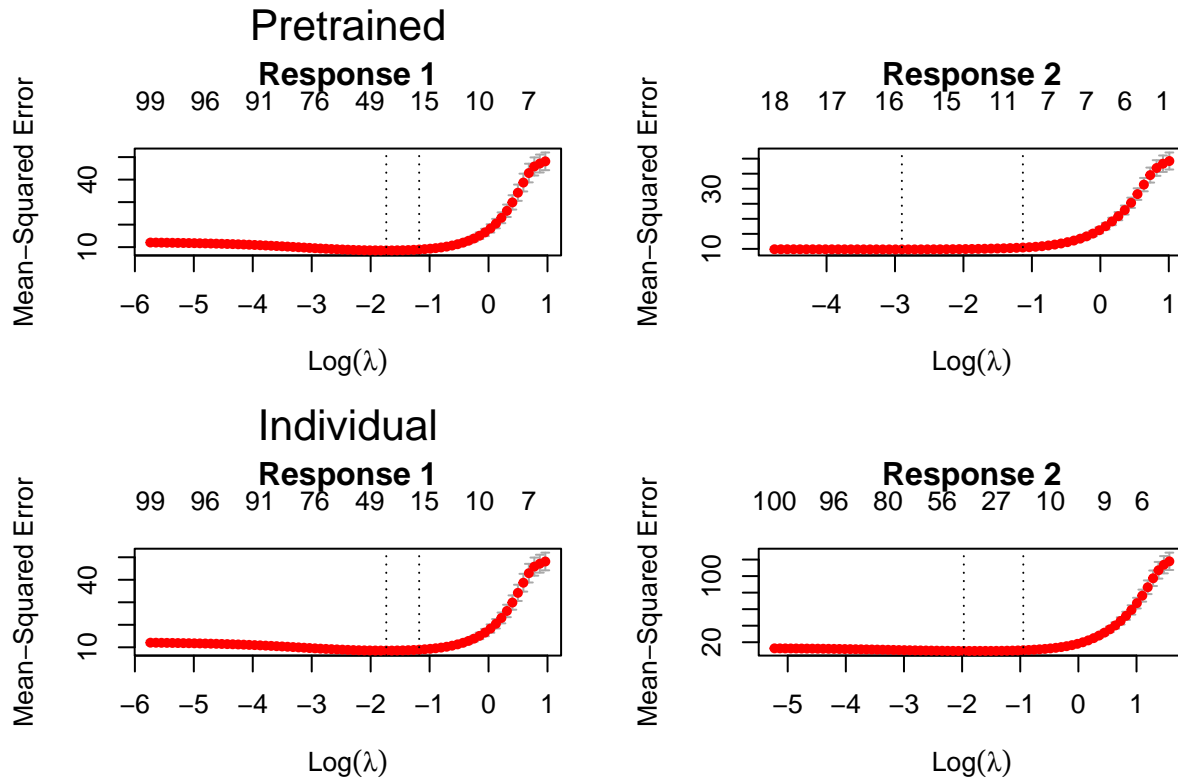
Now, x is a list of length two:

```
str(x)
#> List of 2
#> $ : num [1:300, 1:100] -1.207 0.277 1.084 -2.346 0.429 ...
```

```
#> $ : num [1:300, 1:100] -1.493 0.303 1.172 -2.316 0.224 ...
```


We can call `ptLasso`, `cv.ptLasso`, `plot` and `predict` just as before:

```
fit = ptLasso(x, y, use.case = "timeSeries", alpha = 0)
plot(fit) # Plot the fitted model
```



```
predict(fit, xtest, ytest = ytest) # Predict using the fitted model
#>
#> Call:
#> predict.ptLasso(object = fit, xtest = xtest, ytest = ytest)
#>
#>
#>
#> alpha = 0
#>
#> Performance (Mean squared error):
#>
#>          mean response_1 response_2
#> Pretrain   11.92         12.1      11.75
#> Individual 11.46         12.1      10.82
#>
#> Support size:
#>
#> Pretrain   36 (16 common + 20 individual)
#> Individual 61
#>
# With cross validation:
cvfit = cv.ptLasso(x, y, use.case = "timeSeries")
plot(cvfit, plot.alphahat = TRUE) # Plot cross-validated model
```

Average performance over 2 responses



```
predict(cvfit, xtest, ytest = ytest) # Predict using cross-validated model
#>
#> Call:
#> predict.cv.ptLasso(object = cvfit, xtest = xtest, ytest = ytest)
#>
#>
#>
#> alpha = 0.4
#>
#> Performance (Mean squared error):
#>
#>      mean response_1 response_2
#> Pretrain   15.73      12.11    19.35
#> Individual 11.53      12.11    10.96
#>
#> Support size:
#>
#> Pretrain   54 (19 common + 35 individual)
#> Individual 65
```